



The model achieved 85% accuracy on the validation set.



The dataset I used :

<https://drive.google.com/file/d/1kC60RGO3rcScVk7HY-s7tTMJeMbADfh1/view>

Selected Code:

"""

Code to train the model

```
"""
```

```
import tensorflow as tf
```

```
import numpy as np
```

```
import time
```

```
from datetime import timedelta
```

```
from dataset import load_cached
```

```
#from matplotlib.image import imread
```

```
import cv2,sys, argparse
```

```
#Initializing the conv and max_pool layers
```

```
#####
```

```
def new_conv_layer(input,          # The previous layer.
                   num_input_channels, # Num. channels in prev. layer.
                   filter_size,      # Width and height of each filter.
                   num_filters):     # Number of filters.
```

```
    # Shape of the filter-weights for the convolution.
```

```
    shape = [filter_size, filter_size, num_input_channels, num_filters]
```

```
    # Create new weights aka. filters with the given shape.
```

```
    weights = tf.Variable(tf.truncated_normal(shape, stddev=0.05))
```

```
    # Create new biases, one for each filter.
```

```
    biases = tf.Variable(tf.constant(0.05, shape=[num_filters]))
```

```
    layer = tf.nn.conv2d(input=input,
                          filter=weights,
                          strides=[1, 2, 2, 1],
                          padding='VALID')
```

```

# A bias-value is added to each filter-channel.

layer += biases

return layer

#####

def max_pool(layer, ksize, strides):
    layer = tf.nn.max_pool(value=layer,
                           ksize=ksize,
                           strides = strides,
                           padding = 'VALID')

    return layer

#####

def new_fc_layer(input,          # The previous layer.
                 num_inputs,    # Num. inputs from prev. layer.
                 num_outputs,   # Num. outputs
                 use_relu=True): # Use Rectified Linear Unit (ReLU)?

    # Create new weights and biases.

    weights = tf.Variable(tf.truncated_normal([num_inputs, num_outputs],
                                              stddev=0.05))

    biases = tf.Variable(tf.constant(0.05, shape=[num_outputs]))

    #Include Drop-out as well to avoid overfitting
    #x_drop = tf.nn.dropout(input, keep_prob=keep_prob_input)

    # Calculate the layer as the matrix multiplication of
    # the input and weights, and then add the bias-values.
    layer = tf.matmul(input, weights) + biases

```

```

    # Use ReLU?

    if use_relu:

        layer = tf.nn.relu(layer)

    return layer

#####

def flatten_layer(layer):

    # Get the shape of the input layer.

    layer_shape = layer.get_shape()

    # The shape of the input layer is assumed to be:

    # layer_shape == [num_images, img_height, img_width, num_channels]

    # The number of features is: img_height * img_width * num_channels
    num_features = layer_shape[1:4].num_elements()

    layer_flat = tf.reshape(layer, [-1, num_features])

    # The shape of the flattened layer is now:

    # [num_images, img_height * img_width * num_channels]

    return layer_flat, num_features

#####

class Model:

    def __init__(self, in_dir, save_folder=None):

        dataset = load_cached(cache_path='my_dataset_cache.pkl',
in_dir=in_dir)

        self.num_classes = dataset.num_classes

```

```

        image_paths_train, cls_train, self.labels_train =
dataset.get_training_set()

        image_paths_test, self.cls_test, self.labels_test =
dataset.get_test_set()

#####IMAGE
PARAMETERS#####

        self.img_size = 128

        self.num_channels = 3

        self.train_batch_size = 64

        self.test_batch_size = 64

#####
#####

        self.x = tf.placeholder(tf.float32, shape=[None,
self.img_size,self.img_size,self.num_channels], name='x')

        self.x_image = tf.reshape(self.x, [-1, self.img_size,
self.img_size, self.num_channels])

        self.y_true = tf.placeholder(tf.float32, shape=[None,
self.num_classes], name='y_true')

        self.y_true_cls = tf.argmax(self.y_true, axis=1) #The True class
Value

        self.keep_prob = tf.placeholder(tf.float32)

        self.keep_prob_2 = tf.placeholder(tf.float32)

        self.y_pred_cls = None

        self.train_images= self.load_images(image_paths_train)

        self.test_images= self.load_images(image_paths_test)

        self.save_folder=save_folder

        self.optimizer,self.accuracy = self.define_model()

def load_images(self,image_paths):

```

```

# Load the images from disk.

images = [cv2.imread(path,1) for path in image_paths]

# Convert to a numpy array and return it in the form of
[num_images,size,size,channel]

#print(np.asarray(images[0]).shape)

return np.asarray(images)

def define_model(self):

    #Convolution Layer 1

    filter_size1 = 10          # Convolution filters are 10 x 10
    num_filters1 = 24          # There are 24 of these filters.

    # Convolutional Layer 2

    filter_size2 = 7           # Convolution filters are 7 x 7
    num_filters2 = 48          # There are 48 of these filters.

    # Convolutional Layer 3

    filter_size3 = 11          # Convolution filters are 11 x 11
    num_filters3 = 96          # There are 96 of these filters.

    # Fully-connected layer

    fc_size = 96

    layer_conv1 = new_conv_layer(input=self.x_image,
                                  num_input_channels=self.num_channels,
                                  filter_size=filter_size1,
                                  num_filters=num_filters1)

    #Max Pool Layer

    ksize1 = [1,4,4,1]

    strides1 = [1,2,2,1]

```

```

layer_max_pool1 = max_pool(layer_conv1,ksize1, strides1)

#Convolutional Layer 2
layer_conv2 = new_conv_layer(input=layer_max_pool1,
                              num_input_channels=num_filters1,
                              filter_size=filter_size2,
                              num_filters=num_filters2)

#Max Pool Layer
ksize2 = [1,2,2,1]
strides2 = [1,1,1,1]
layer_max_pool2 = max_pool(layer_conv2,ksize2, strides2)

#Convolutional Layer 3
layer_conv3 = new_conv_layer(input=layer_max_pool2,
                              num_input_channels=num_filters2,
                              filter_size=filter_size3,
                              num_filters=num_filters3)

#Flatten
layer_flat, num_features = flatten_layer(layer_conv3)

#Relu Layer
layer_relu = tf.nn.relu(layer_flat)

#Fully-Connected Layer1
layer_fc1 = new_fc_layer(input=layer_relu,
                          num_inputs=num_features,
                          num_outputs=fc_size,
                          use_relu=True)

#Fully-Connected Layer2
layer_fc2 = new_fc_layer(input=layer_fc1,
                          num_inputs=fc_size,

```

```

        num_outputs=self.num_classes,
        use_relu=False)

    #Predict the class
    y_pred = tf.nn.softmax(layer_fc2)
    self.y_pred_cls = tf.argmax(y_pred,
dimension=1,name="predictions")

    #Cost Function
    cross_entropy =
tf.nn.softmax_cross_entropy_with_logits(logits=layer_fc2,
labels=self.y_true)
    cost = tf.reduce_mean(cross_entropy)
    optimizer = tf.train.AdamOptimizer(learning_rate=1e-
4).minimize(cost)

    #Predict
    correct_prediction = tf.equal(self.y_pred_cls, self.y_true_cls)
    accuracy = tf.reduce_mean(tf.cast(correct_prediction,
tf.float32))

    return optimizer, accuracy

def random_batch(self):
    # Number of images in the training-set.
    num_images = len(self.train_images)

    # Create a random index.
    idx = np.random.choice(num_images,
        size=self.train_batch_size,
        replace=False)

```



```

        # Use the random index to select random x and y-values.
        x_batch = self.train_images[idx]
        y_batch = self.labels_train[idx]

    return x_batch, y_batch

def print_test_accuracy(self, sess):

    # Number of images in the test-set.
    num_test = len(self.test_images)

    # Allocate an array for the predicted classes which
    # will be calculated in batches and filled into this array.
    cls_pred = np.zeros(shape=num_test, dtype=np.int)

    i = 0

    while i < num_test:
        # The ending index for the next batch is denoted j.
        j = min(i + self.test_batch_size, num_test)

        images = self.test_images[i:j]

        labels = self.labels_test[i:j]

        # Create a feed-dict with these images and labels.
        feed_dict = {self.x: images,
                     self.y_true: labels,
                     self.keep_prob: 1,
                     self.keep_prob: 1}

```

```

        cls_pred[i:j] = sess.run(self.y_pred_cls,
feed_dict=feed_dict)

        # Set the start-index for the next batch to the
        # end-index of the current batch.
        i = j

    # Create a boolean array whether each image is correctly
    classified.

    correct = (self.cls_test == cls_pred)

    # Classification accuracy is the number of correctly classified
    # images divided by the total number of images in the test-set.
    acc = float(correct.sum()) / num_test

    # Print the accuracy.
    msg = "Accuracy on Test-Set: {0:.1%} ({1} / {2})"
    print(msg.format(acc, correct.sum(), num_test))

def optimize(self, num_iterations):
    # Ensure we update the global variable rather than a local copy.
    global total_iterations
    total_iterations = 0
    saver = tf.train.Saver()

    # Start-time used for printing time-usage below.
    start_time = time.time()

    with tf.Session() as sess:
        #global_step_int = tf.train.get_global_step(sess.graph)
        sess.run(tf.global_variables_initializer())

```

```

for i in range(total_iterations,
               total_iterations + num_iterations):

    # Get a batch of training examples.
    # x_batch now holds a batch of images and
    # y_true_batch are the true labels for those images.
    x_batch, y_true_batch = self.random_batch()

    feed_dict_train = {self.x: x_batch,
                       self.y_true: y_true_batch}
                       #self.keep_prob: 0.5,
                       #self.keep_prob: 0.5}

    sess.run([self.optimizer], feed_dict=feed_dict_train)

    # Print status every 100 iterations.
    if i % 100 == 0:
        # Calculate the accuracy on the training-set.
        feed_dict_acc = {self.x: x_batch,
                        self.y_true: y_true_batch}
                        #self.keep_prob: 1,
                        #self.keep_prob: 1}

        acc = sess.run(self.accuracy,
                        feed_dict=feed_dict_acc)

        # Message for printing.
        msg = "Optimization Iteration: {0:>6}, Training
Accuracy: {1:>6.1%}"

        # Print it.

```

```

        print(msg.format(i + 1, acc))

        # Update the total number of iterations performed.
        total_iterations += num_iterations

        # Ending time.
        end_time = time.time()

    if i%100 ==0:
        #Calculate the accuracy on the test set every 100
iterations
        self.print_test_accuracy(sess)

    if i%500 == 0:
        #Saves every 500 iterations
        saver.save(sess,
os.path.join(self.save_folder,'model')) #Change this according to your
convenience

        # Difference between start and end-times.
        time_dif = end_time - start_time
        self.print_test_accuracy(sess)
        # Print the time-usage.
        print("Time usage: " +
str(timedelta(seconds=int(round(time_dif)))))
        saver.save(sess,
os.path.join(self.save_folder,'model_complete'))

def parse_arguments():
    parser = argparse.ArgumentParser(description='Training Network')

```

```
    parser.add_argument('--
in_dir',dest='in_dir',type=str,default='cracky')
    parser.add_argument('--
iter',dest='num_iterations',type=int,default=1500)
    parser.add_argument('--
save_folder',dest='save_folder',type=str,default=os.getcwd())
    return parser.parse_args()

def main(args):
    args=parse_arguments()
    num_iterations = args.num_iterations

    model = Model(args.in_dir,args.save_folder)
    model.optimize(num_iterations)

if __name__ == '__main__':
    main(sys.argv)
```